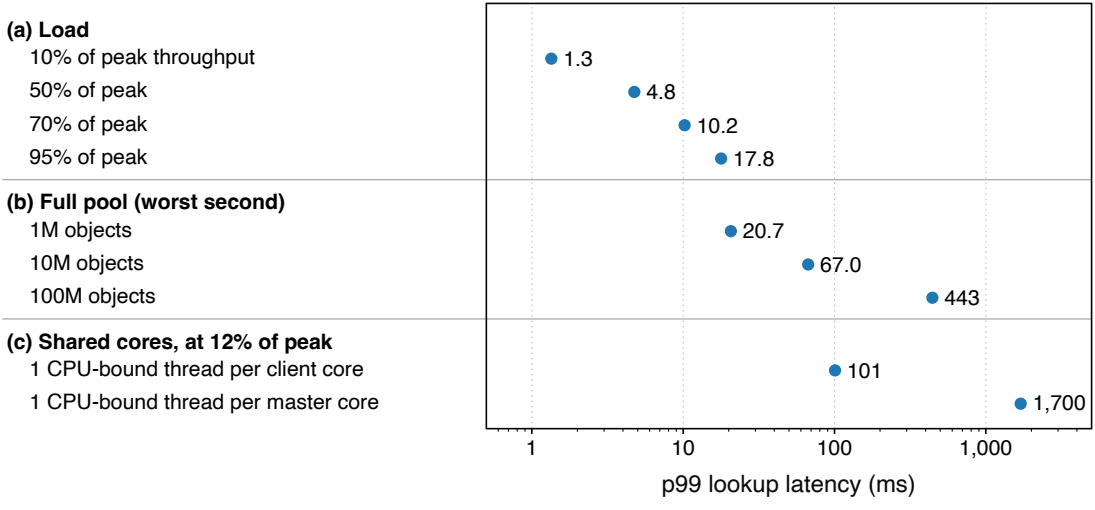
(a) Load
10% of peak throughput
50% of peak
70% of peak
95% of peak
1.3
4.8
10.2
17.8
(b) Full pool (worst second)
1M objects
10M objects
100M objects
20.7
67.0
443
(c) Shared cores, at 12% of peak
1 CPU-bound thread per client core
1 CPU-bound thread per master core
101
1,700
1
10
100
1,000
p99 lookup latency (ms)