How much work does serving send the master?

Demand from one Kimi cluster

56K prompt tokens/s

Mooncake's released 1-hour trace, replayed on 16 nodes × 8 GPUs

↓ one key per 64-token page per GPU: ÷ 64 × 8

7.0K lookups/s

↓ only the 41% of pages not already cached are written

2.9K writes/s

What one master handles

5.1M lookups/s

175K writes/s

measured in E1: 16 cores, no network

Serving uses about **0.1%** of the master's lookup capacity and **2%** of its write capacity. Request rate is not the problem.

Tool & agent trace (Mooncake, FAST '25): 23,608 requests in 1 hour, 8,596 prompt tokens each on average, 59% cached; the conversation trace gives 40K tokens/s. Hourly averages; peaks are higher. Page layout as in SGLang's HiCache.